

Charybdis Manual for Metabarcoding

Evan Krell, Martin French, Sharon Magnuson, Chris Bird

Genomics Core Lab

Texas A&M University - Corpus Christi

ABSTRACT

1 NOTICE

This manual was written for the Texas A&M University - Corpus Christi HPC. This system uses SLURM for job management. The path references are for this system, but an experienced Linux user should be able to set it up for their own system. A configuration file is currently being setup that will make it easier for users to run. Workstations should be able to run it without issue, so long as the dependencies are met.

2 INTRODUCTION

Charybdis involves numerous scripts in a variety of programming languages. A charybdis project is done by creating a pipeline of these scripts and setting the parameters of each. Because the needs of individual projects differs, it is difficult to create a one-size-fits-all program. Therefore, the use of small, modular programs appears to be the most effective. This manual describes each script in the charybdis pipeline, demonstrates an example pipeline, and offers guidance on setting up databases and running on an HPC.

Multiple assignment methods are supported and may be run simultaneously. For example, a user can supply options to use both BLAST and VSEARCH. At the assignment step, the pipeline splits and both methods are run in parallel. Each method produces a number of resulting files. These files are differentiated by including the assignment method used in the path name.

Taxonomic assignment is not required. If no assignment databases are specified, the pipeline can still be used for just the filtering and clustering of reads.

2.1 Charybdis

All of the code resides in a repository called `charybdis`. The repository is located on the HPC at `/work/hobi/GCL/charybdis`. The name resulted from finding a critter from species *Charybdis* from our first set of metabarcoding results.

3 PIPELINE PARAMETERS

You need a number of files and bits of information in order to run the pipeline. The following describes each parameter, and (if relevant), how to obtain them. See section 5 for actual examples of these parameters used in a project.

3.0.1 Project Name (-p)

The name of the project is used as a prefix to many files generated throughout the pipeline. Additionally, the input files need to use this project name in their filenames.

3.0.2 Input Directory (-i)

Certain files need to be placed in this directory before executing the pipeline. Instead of forcing the user to specify several files, they are placed in a single location. The pipeline finds files by using the project name specified with the `-p` option. The required files for running the generic pipeline script is shown in Table 1.

3.0.3 Output Directory (-o)

Files generated by pipeline will be placed in here.

Table 1. Files required to exist in the input directory.
The variable project name is represented by <projectname>.

<projectname>_forward.fastq	FASTQ with forward reads
<projectname>_reverse.fastq	FASTQ with forward reads
<projectname>.barcodes.txt	Tab-delimited file to match up barcodes to samples
<projectname>.sampledescs.csv	CSV that matches descriptive names to samples

41 **3.0.4 Chunks (-n)**

42 Number of parallel threads to spawn. Use all the cores available if running on an HPC node. But if
43 running on desktop, use less than your number of available cores. For the 40-core workstation in Bird's
44 computer room, I would go with 35.

45 **3.0.5 BLAST Database (-b)**

46 Any valid BLAST database. We typically uses the the nucleotide (nt) database from NCBI.
47 `ftp://ftp.ncbi.nlm.nih.gov/blast/db/`
48 See 8.1 for information on obtaining a local BLAST database.

49 **3.0.6 BLAST Ignore File (-d)**

50 A text file containing GIs of sequences in the BLAST database to ignore. It is fine to have GIs that are
51 not actually in the BLAST database used. Good for filtering environmental DNA and such. See 8.1.1 for
52 information on obtaining a local BLAST database.

```
53 [ekrell@hpcm Simons-H1]$ head -n 5 NCBI.NT-AUG2018.env.gi
54     1000001401
55     1000001403
56     1000001405
57     1000001407
58     1000001409
```

59 **3.0.7 VSEARCH Database (-v)**

60 VSEARCH is an alternative taxonomic assignment program Rognes et al. (2016). A valid VSEARCH
61 database is a FASTA file. But because FASTA files are a loose standard, not all are compatible. VSEARCH
62 gets unhappy regarding certain formats of the sequence header. The pipeline supports two formats of
63 FASTA. One uses data obtained from Barcode of Life (Ratnasingham and Hebert, 2007) and another from
64 NCBI. See 8.2 for information on obtaining a local VSEARCH database.

65 An example of a valid FASTA entry is as follows (data is from (Ratnasingham and Hebert, 2007)).
66 Note that the | symbol is used to separate metadata field in the header. The first field is a unique sequence
67 ID. The second is a scientific name. The third is the genetic region of the sequence. The fourth is the
68 Barcode of Life Database (BOLD) sequence ID. It could be a NCBI GI or Accession instead.

```
69     >GBMAA467X14|Pomphorhynchus.laevis|COIX5P|KF559289
70     ATGTATGTTTGGTTGGTGTGTGAGGGGGGCTAATGGGGTTTCTATAAGACTATTAATTCGA
```

71 The script `/work/hobi/GCL/charybdis/bin/vsearch.getTAXIDfromBOLDseqid.sh`
72 will get the NCBI taxonomic ID from a BOLD sequence ID.

73 **3.0.8 Chimera Database (-c)**

74 FASTA database used for detection on chimeric reads. Chimera detection can be done using a database or
75 using the *de novo* method Rognes et al. (2016). Currently the pipeline is hard-coded to use the *de novo*
76 method, but still requires a chimera database that will not be used. This issue is described in section 6.3.

77 **3.0.9 Target Sequence Length (-x)**

78 The target sequence length is the length of a read expected using PCR. The forward and reverse primers
79 extract a portion of DNA between them. The value of this option is based on the primers you used, so
80 check their documentation or ask whoever did the PCR.

81 **3.0.10 Charybdis Scripts Directory (-g)**

82 Instead of setting the environment `$PATH` variable or placing the scripts in a systemwide bin directory,
83 the directory holding all the pipeline scripts is an option. Less conventional than `$PATH`-based methods,
84 but very easy to deal with.

85 **3.0.11 NCBI Taxonomy Database (-t)**

86 This database is used to get an organism's scientific name from a numeric taxonomic ID. Taxonomy
87 database downloaded from `ftp://ftp.ncbi.nlm.nih.gov/pub/taxonomy/`.

88 It is the same data as `https://www.ncbi.nlm.nih.gov/Taxonomy/Browser/wwwtax.cgi`

89 **4 PIPELINE ARCHITECTURE**

90 Pipelines are bash scripts that just string together a group of other scripts/programs to perform the metabar-
91 coding process. Example scripts are included in `/work/hobi/GCL/charybdis/pipelines`. The
92 script called `charybdis_generic` is expected to handle most projects with little customization.

93 The section will describe all of the steps of the `charybdis_generic` pipeline. Hopefully others
94 will be able to use this information to create a pipeline tailored to other metabarcoding projects.

95 Throughout the text, *attribute* refers to a piece of metadata included in a FASTA sequence header.
96 These are in the form of key-value pairs. Many tools in the OBITOOLS package create and modify
97 attributes, so we made many of our own scripts follow this convention.

98 **4.1 Merge Reads**

99 **Script:** `/work/hobi/GCL/charybdis/bin/mergeReads.sh`

100 **Slurm wrapper:** `/work/hobi/GCL/charybdis/bin/mergeReads.slurm`

101 The forward and reverse reads are merged into a single FASTA file, and basic quality filtering is done.
102 The reads are matched to sample ID which are included as metadata in the FASTA sequence headers.

103 **4.1.1 Split FASTQ files**

104 In order to run the next sections in parallel, the forward and reverse read FASTQ files are split into equally
105 sized chunks using `fastq-splitter.pl`.

106 **4.1.2 Join forward and reverse reads**

107 After this point, the pipeline uses a single file of merged forward and reverse reads. This is done using
108 OBITOOL's `illuminapairedend` tool. We have it reject reads whose alignment score is under a
109 threshold of 40.

110 **4.1.3 Convert FASTQ to FASTA**

111 Convert the merged files to FASTA. This is done because many tools operate only on FASTA and because
112 we don't need the sequencer quality information after this point. We can always use the unique sequence
113 ID to access it from the origin FASTQ, if needed.

114 **4.1.4 Remove unaligned sequences**

115 We use OBITOOL's `obigrep` to select only those sequences whose alignment mode is "joined". `Obigrep`
116 can search for patterns of values of specified attributes.

117 **4.1.5 Remove sequences whose alignment length is below threshold**

118 We use OBITOOL's `obigrep` to select only those sequences whose alignment length is \geq a threshold. We
119 have this threshold set to 20.

120 **4.1.6 Match sequences to sample ID based on barcode**

121 OBITOOL's `ngsfilter` is used to attach an attribute called *sample* to each sequence. The match is allowed
122 to have 2 sequence errors by setting the flag `-e 2`.

123 **4.1.7 Concatenate into single FASTA**

124 The small chunks are concatenated into a single FASTA file.

125 4.2 Filter Reads

126 **Script:** /work/hobi/GCL/charybdis/bin/ObiToolsPipeline.sh

127 **Slurm wrapper:** /work/hobi/GCL/charybdis/bin/filterReads.slurm

128 This was the first script written. At the time, the size of the project was not known and this file was
129 intended to be the entire pipeline. Thus, the script is called `ObiToolsPipeline.sh` instead of the
130 more appropriate `filterReads.sh`.

131 4.2.1 Split into smaller files by sample ID

132 As before, we form a smaller number of files for parallel processing. However, the sample ID is used to
133 split files instead of an arbitrary size. The sample ID is added to the FASTA filename to keep track.

134 4.2.2 Remove unneeded attributes

135 OBITOOLS added a huge number of attributes in a previous step. We get rid of the irrelevant ones with
136 `obiannotate`. Note that attributes are never truly discarded since the sequence ID can always be used
137 to query the intermediate FASTAs generated at each step of the pipeline.

138 4.2.3 Keep only unique sequences

139 In order to lower the filesize (often substantially), duplicated sequences are merged into a single represen-
140 tative sequence. The `obimerge` attribute is used to record how many sequences are represented by that
141 sequence. This is done using `obiuniq`.

142 4.2.4 Remove PCR errors

143 PCR errors are filtered using `obiclean`. After numerous experiments, we found that the following
144 options were performing well for a set of COI sequences from Gulf of Mexico fish. Results from projects
145 where we reused these settings have appeared to be fine.

Table 2. Settings for `obiclean`

-r	0.5	Controls when less abundant sequences are labeled as variants of a similar more abundant
-d	1	Max number of differences to be considered variants
-H		Keep only a representative of a set of variants.

146 4.2.5 Remove low read length

147 The amplified sequence should be a certain size. Using our COI primers, we expect the size to be 313
148 base pairs. However, the actual size is often a bit lower or higher. You supply your base pairs length as a
149 pipeline argument. The allowed range is that size ± 15 .

150 4.2.6 Remove chimeras

151 VSEARCH is used to remove chimeras using the *de novo* method.

152 4.2.7 Concatenate into single FASTA

153 The small chunks are concatenated into a single FASTA file.

154 4.3 Collapse into OTUs

155 **Script:** /work/hobi/GCL/charybdis/bin/ClusterOTU.sh

156 **Slurm wrapper:** /work/hobi/GCL/charybdis/bin/ClusterOTU.slurm

157 4.3.1 Cluster with CROP

158 CROP clusters sequences into OTUs (Hao et al., 2011). CROP has a number of parameters that control
159 the OTU assignment. Assigning these parameters is tricky, but the CROP manual has a recommended
160 process for determining the values (TingChenLab, 2017). We initially used the recommendation, but then
161 modified it to get better results. **However, we don't currently remember how we figured out some of the**
162 **values used in our process.** You can check the manual and compare to ours, but note that `-b` and `-z` should
163 only affect performance, not clustering. However, we set `-e` based on `-z` (following the manual), which
164 does affect results. **We need to investigate further.**

Table 3. Settings for CROP.

NUMSEQS is the total number of sequences. *BP* is the target sequence length.

-r	5	Clusters of size <i>leqr</i> are considered rare
-s		Specifies that 97% similarity needed to cluster
-b	$NUMSEQS/50$	Number of blocks
-z	$(150000/BP) - 0.1 * (150000/BP)$	Max number of sequence a block can hold
-e	$[-z] * 10$	Number of MCMC iterations

4.3.2 Correct the OTU size

CROP records what sequences end up in an OTU, as well as the total number of sequences represented by an OTU. However, CROP has no knowledge that obiclean and obimerge already collapsed duplicate/similar sequences into. We call `CROP_size_fix.sh` to get the actual OTU counts.

4.4 Taxonomic Assignment

Taxonomic assignment can be done using either BLAST or VSEARCH. (SAP coming soon). The exact scripts and slurm wrappers depend on method selected.

4.4.1 Run taxonomic assignment program in parallel

Split sequences into chunks and BLAST in parallel. Uses the BLAST ignore list (provided in parameters) to skip assignment to unwanted sequences. For example, we typically try to avoid assignment to environmental sequences.

4.4.2 Concatenate into single FASTA

The small chunks are concatenated into a single FASTA file.

4.4.3 Convert to Charon format

Charon is the name of our own CSV format for assigned sequences. The purpose is that we can write conversion scripts to have assignments from various taxonomic assignment programs into a standard. The Charon file has columns specified by Table 4.

Table 4. Columns for Charon CSV file

Query Sequence ID	Assignment Sequence ID	Scientific Name	NCBI GI	Number of sequences in OTU
-------------------	------------------------	-----------------	---------	----------------------------

4.5 OTUs VS Samples

Formerly: *Critters VS Tubes*.

Will explain soon, but basically just a script that makes a column for each sample and a row for each OTU. The cells have a count of the number of sequences of that OTU in that sample. Also has columns of the scientific names for various phylogenetic levels. Also has extra information columns such as the sequence, BLAST scores, etc. No limit on added extra columns to this CSV file.

5 EXAMPLE RUN

5.1 Prepare Data

5.2 Run Pipeline Script

The following bash script executes `charybdis_generic.sh`, an implementation of the GCL pipeline. The code segment shown is a single line of bash, using `\` for clarity to break it into one line per option.

```
bash /work/hobi/GCL/charybdis/pipelines/charybdis_generic.sh \
-p Simons-H1.ATTACTCG-TAATCTTA.L001 \
-i /work/hobi/GCL/20180330-Simons-Simons-H1.in \
-o /work/hobi/GCL/20180330-Simons-Simons-H1.out \
-n 20 \
-b /work/hobi/GCL/db/NCBI.BLAST.DBs/nt \
-v /work/hobi/GCL/db/BOLD.FULL.fasta \
```



```

260         08-A, Larimus_fasciatus
261     -o Output directory.
262     Files generated by the pipeline, including numerous intermediate file, are placed here.
263     -n Number of parallel instances.
264     The nodes on TAMUCC's HPC cave 20 cores, so 20 were used.
265     -b BLAST database.
266     Specify path to BLAST database. We are using the nucleotide (nt) database,
267     which is downloaded from ftp://ftp.ncbi.nlm.nih.gov/blast/db/.
268     -c VSEARCH database
269     Specify path to VSEARCH database. We use a FASTA file generated from the BOLD database.
270     -d List of GIs for BLAST to ignore This is a list of GIs, where each GI specifies a sequence in the BLAST
271     database to ignore. This file is mandatory, but may be blank (See Section 6.2).
272     We attempt to remove environmental DNA, so our list has GIs of environmental sequences.
273     [ekrell@hpcm Simons-H1]$ head -n 5 ...../environmental.NCBI_nucl__SORTED.gi
274         1000014437
275         1000014439
276         1000014441
277         1000014443
278         1000014445
279     -c Chimera Database. We use a FASTA file generated from the BOLD database. But, we ended up doing de novo
280     chimera detection. Currently this file is required and just not being used (See section 6.3).
281     -x Target sequence length (basepairs). This number is based on the primers used during PCR.
282     -g Directory of charybdis scripts.
283     ls /work/hobi/GCL/charybdis/bin
284     \
285     AddBlast.slurm
286     AddVsearch.slurm
287     AssignToMarkers.sh
288     AssignToMarkers.slurm
289     blast_10custom_to_charon_OTU.R
290     BlastMeta.sh
291     BlastMeta.slurm
292     ClusterOTU.sh
293     ClusterOTU.slurm
294     CombineAndCROP.sh
295     crittersVStubes_OTU.R
296     CROP_size_fix.sh
297     CROP.slurm
298     determine_marker.R
299     determine_marker.sh
300     fastq-splitter.pl
301     filterReads.slurm
302     mergeReads.sh
303     mergeReads.slurm
304     ObiToolsPipeline.sh
305     ObiToolsPipeline.slurm
306     ObiToolsPipeline_stats.sh
307     OTU_CVT_addBlast.R
308     OTUvsTube.slurm
309     split_by_marker.R
310     vsearch_blast6custom_to_charon-BLASTDB_OTU.R
311     vsearch_blast6custom_to_charon_OTU.R
312     vsearch_getTAXIDfromBOLDseqid.sh
313     Vsearch.sh
314     Vsearch.slurm
315     -t Directory with NCBI taxonomy database Taxonomy database downloaded from ftp://ftp.ncbi.nlm.
316     nih.gov/pub/taxonomy/.
317     -g Directory of charybdis scripts.
318     [ekrell@hpcm pipelines]$ ls /work/hobi/GCL/db/TAXO
319     Accession2Taxid delnodes.dmp gc.prt merged.dmp nodes.dmp taxdump.tar.gz
320     citations.dmp division.dmp gencode.dmp names.dmp readme.txt

```

306 6 AREAS FOR IMPROVEMENT

307 6.1 Support NCBI Accession IDs

308 NCBI is phasing out GIs (sequence IDs), which this pipeline uses at various steps (NCBI, 2016). Accession IDs are
309 the current standard.

310 6.2 Pipeline should not require BLAST ignore list

311 Currently this is required. If you don't need it, you can supply a blank file. Not very elegant. Would be simple to
312 make it optional.

313 6.3 Chimera database should not be mandatory

314 The generic pipeline uses *de novo* chimera detection, but still requires the chimera database. A better solution would
315 be to default to *de novo*, but use the database if supplied.

316 7 STATISTICAL ASSIGNMENT PACKAGE (SAP)

317 **We have SAP working, but not meshed into the pipeline yet. So it has its own section, at least for now.**

318 SAP (Munch et al., 2008) differs from typical assignment software (BLAST, VSEARCH) in that it is based on
319 computing a posterior probability that the query sequence is a member of a particular clade. Two major phases are
320 involved in an SAP taxonomic assignment. First, BLAST is executed to find a set of similar sequences. Since similar
321 sequences should indicate similar biological properties, the set of accepted BLAST sequences are called the set of
322 homologues with the query. To strengthen the assumption that the sequences are homologous, there is a limit on
323 how distant the furthest apart homologues can be. Several parameters are used to control BLAST's sequence search
324 and also to control how SAP chooses which of those sequences to include in the homologue set. It can and does
325 happen that a single species has several sequences in the database. However, a particular taxonomic group appears
326 only once in the phylogenetic tree. Therefore, SAP focuses on diversity and only uses the best match for a given
327 species. While this is necessary to construct meaningful trees, it also means that the species which only appears once
328 in a set of BLAST results is given equal weight to a predominate species result. The set of homologues and the query
329 sequence are used to sample a large number of trees using Markov Chain Monte Carlo. The posterior probability for
330 any taxonomic group is based on the proportion of trees in which the query sequence was in that group. A simple
331 case that shows how this can differ from BLAST occurs when several species from the same genus are present in the
332 homologue set. Where BLAST would simply list the hits in order of sequence ID, SAP determines that it cannot
333 reliably be determined what species the query belongs to. A top hit may differ little from the hits just below it, and it
334 is essentially arbitrary to assume that the top hit is the exact species when it scored closely against other species as
335 well. The SAP result would give a low posterior probability to any particular sequence, but the genus would have a
336 very high probability.

337 *The above text taken from our work-in-progress manuscript "An Evaluation of Software for Taxonomic Assignment*
338 *with DNA Metabarcoding"*

339 7.1 Fixing SAP

340 7.2 Running SAP

341 7.3 Evaluating SAP

342 7.3.1 Categorize SAP results

343 7.3.2 Generate SAP report

344 7.4 SAP outstanding issues

345 7.5 Complete SAP example

346 7.6 SAP outstanding issues

347 8 DATABASE CREATION

348 8.1 BLAST Database

349 The pipeline assumes that you are using the NCBI provided BLAST databases. While the BLAST component will
350 work for any BLAST database, further steps in the pipeline expect particular formatting of results. This is based on
351 what information is available in the BLAST results. Any filters may be applied to this database, as long as the format
352 is maintained.

353 *NCBI BLAST Database:*

354 `ftp://ftp.ncbi.nlm.nih.gov/blast/db`

355 *Instructions for download and use:*

356 `ftp://ftp.ncbi.nlm.nih.gov/blast/documents/blastdb.html`

357 8.1.1 Filtering BLAST database

358 The NCBI contains a large amount of environmental samples and similar junk. We want to remove this material.

359 How to filter environmental samples.

- 360 • Get list of unwanted entries. Use an NCBI Entrez query.

361 Query:

362 `"environmental_samples"[organism] OR metagenomes[orgn] OR sp[Title]`

363 URL: `https://www.ncbi.nlm.nih.gov/nuccore/?term=%22environmental+samples%22%5Borganism%5D+OR+metagenomes%5Borgn%5D+OR+sp%5BTitle%5D`

- 365 • At the above webpage, locate the number of entries found by the query. For example, in the line `Items : 1`
366 `to 20 of 21247214`. You want to copy the number 21247214 (or your actual number).

- This number will replace <number> in the below command.
- ```
bash /work/hobi/GCL/charybdis/bin/getEnvGIlist.sh <number> > \
/work/hobi/GCL/db/NCBI_NT_SEPT2018.env.gi
```
- There are two ways to use the downloaded GI list. You can use it each time you run the pipeline, using the `-d` (BLAST ignore file) option. These items will be ignored when searching BLAST. It can also be used with create an alias BLAST database. After creating this with `blastdb_aliastool`, you can use the alias database for the `-b` option. The alias is really a middle-man between your search and the full BLAST database. It still goes to original database, but with the filter applied. This means your alias database will fail if you move or remove the original BLAST database used to create it.

## 8.2 VSEARCH Database

VSEARCH databases are FASTA files. Because the pipeline looks for information stored in the header, arbitrary FASTAs are not permitted. Currently, the pipeline supports two schemes. One scheme comes from the Barcode of Life (BOLD) database and other by converting the NCBI-style BLAST database to a compatible FASTA.

### 8.2.1 VSEARCH BOLD database

*BOLD FASTA sequence header format*

```
>GBMAA467X14|Pomphorhynchus.laevis|COIX5P|KF559289
ATGTATGTTTTGGTTGGTGTGTGAGGGGGCTAATGGGGTTTCTATAAGACTATTAATTCTGA
```

**Need steps on how to get BOLD database. Unless we instead just deprecate using BOLD...**

### 8.2.2 VSEARCH NCBI database

*NCBI FASTA sequence header format*

```
>X17276.1|Giant_Panda_satellite_1_DNA
GATCCTCCCCAGGCCCTACACCCAATGTGGAACCGGGTCCCGAATGAAAATGCTGCTGTTC
```

How to generate the above from existing NCBI BLAST database (see 8.1).

- (Optional). Use environmental GI list to create filtered BLAST database.

- Convert BLAST database to FASTA

```
blastdbcmd -entry all -db <database> -out <fasta outfile>
```

Example:

```
blastdbcmd -entry all -db /work/hobi/GCL/db/NCBI_BLAST_DBs/nt \
-out /work/hobi/GCL/db/NCBI_BLAST_DBs/NCBI_NT_SEPT2018.fasta
```

- Format FASTA to be VSEARCH compatible

```
sed -e 's/ /|/' -e 's/_/_/' <fasta outfile> \
<vsearchable fasta outfile>
```

Example:

```
sed -e 's/ /|/' -e 's/_/_/' NCBI_NT_DEC2017.fasta > \
NCBI_NT_DEC2017.vsearchable.fasta
```

## REFERENCES

- Hao, X., Jiang, R., and Chen, T. (2011). Clustering 16s rna for otu prediction: a method of unsupervised bayesian clustering. *v*, 27:611–8.
- Munch, K., Boomsma, W., Huelsenbeck, J. P., Willerslev, E., and Nielsen, R. (2008). Statistical assignment of dna sequences using bayesian phylogenetics. *Systematic biology*, 57(5):750–757.
- NCBI (2016). Ncbi is phasing out sequence gis – here’s what you need to know.
- Ratnasingham, S. and Hebert, P. D. (2007). bold: The Barcode of Life Data System (<http://www.barcodinglife.org>). *Mol. Ecol. Notes*, 7(3):355–364.
- Rognes, T., Flouri, T., Nichols, B., Quince, C., and Mahe, F. (2016). VSEARCH: a versatile open source tool for metagenomics. *PeerJ*, 4:e2584.
- TingChenLab (2017). Crop. <https://github.com/tingchenlab/CROP>.