

Texas A&M University - Corpus Christi

Genomics Core Lab

Metabarcoding Manual

Evan Krell, Chris Bird, Martin French, Sharon Magnuson

ABSTRACT

1 INTRODUCTION

The GCL metabarcoding process involves numerous scripts in a variety of programming languages. A metabarcoding project is done by creating a pipeline of these scripts and setting the parameters of each. Because the needs of individual projects differs, it is difficult to create a one-size-fits-all program. Therefore, the use of small, modular programs appears to be the most effective. This manual describes each script in the GCL pipeline, demonstrates an example pipeline, and offers guidance on setting up databases and running on an HPC.

Multiple assignment methods are supported and may be run simultaneously. For example, a user can supply options to use both BLAST and VSEARCH. At the assignment step, the pipeline splits and both methods are run in parallel. Each method produces a number of resulting files. These files are differentiated by including the assignment method used in the path name.

Taxonomic assignment is not required. If no assignment databases are specified, the pipeline can still be used for just the filtering and clustering of reads.

1.1 Charybdis

All of the code resides in a repository called GCL_Charybdis. The repository is located on the HPC at /work/hobi/GCL/GCL_Charybdis. The name resulted from finding a critter from species *Charybdis* from our first set of metabarcoding results.

2 PIPELINE PARAMETERS

You need a number of files and bits of information in order to run the pipeline. The following describes each parameter, and (if relevant), how to obtain them. See section 4 for actual examples of these parameters used in a project.

2.0.1 Project Name (-p)

The name of the project is used as a prefix to many files generated throughout the pipeline. Additionally, the input files need to use this project name in their filenames.

2.0.2 Input Directory (-i)

Certain files need to be placed in this directory before executing the pipeline. Instead of forcing the user to specify several files, they are placed in a single location. The pipeline finds files by using the project name specified with the -p option. The required files for running the generic pipeline script is shown in Table 1.

2.0.3 Output Directory (-o)

Files generated by pipeline will be placed in here.

2.0.4 Chunks (-n)

Number of parallel threads to spawn. Use all the cores available if running on an HPC node. But if running on desktop, use less than your number of available cores. For the 40-core workstation in Bird's computer room, I would go with 35.

Table 1. Files required to exist in the input directory.
The variable project name is represented by <projectname>.

<projectname>_forward.fastq	FASTQ with forward reads
<projectname>_reverse.fastq	FASTQ with forward reads
<projectname>.barcodes.txt	Tab-delimited file to match up barcodes to samples
<projectname>.sampledescs.csv	CSV that matches descriptive names to samples

41 **2.0.5 BLAST Database (-b)**

42 Any valid BLAST database. GCL typically uses the the nucleotide (nt) database from NCBI.
43 `ftp://ftp.ncbi.nlm.nih.gov/blast/db/`

44 **2.0.6 BLAST Ignore File (-d)**

45 A text file containing GIs of sequences in the BLAST database to ignore. It is fine to have GIs that are not
46 actually in the BLAST database used. Good for filtering environmental DNA and such.

47 **2.0.7 VSEARCH Database (-v)**

48 VSEARCH is an alternative taxonomic assignment program Rognes et al. (2016). A valid VSEARCH
49 database is a FASTA file. But because FASTA files are a loose standard, not all are compatible. VSEARCH
50 gets unhappy regarding certain formats of the sequence header. An example of a valid FASTA entry is
51 as follows (data is from (Ratnasingham and Hebert, 2007)). Note that the | symbol is used to separate
52 metadata field in the header. The first field is a unique sequence ID. The second is a scientific name. The
53 third is the genetic region of the sequence. The fourth is the Barcode of Life Database (BOLD) sequence
54 ID. It could be a NCBI GI or Accession instead.

55 `>GBMAA467X14| Pomphorhynchus laevis | COIX5P | KF559289`
56 `ATGTATGTTTGGTGGTGTGTGAGGGGGGCTAATGGGGTTTCTATAAGACTATTAATTCGA`

57 The script `/work/hobi/GCL/GCL_Charybdis/bin/vsearch_getTAXIDfromBOLDseqid.sh`
58 will get the NCBI taxonomic ID from a BOLD sequence ID.

59 **2.0.8 Chimera Database (-c)**

60 FASTA database used for detection on chimeric reads. Chimera detection can be done using a database or
61 using the *de novo* method Rognes et al. (2016). Currently the pipeline is hard-coded to use the *de novo*
62 method, but still requires a chimera database that will not be used. This issue is described in section 5.3.

63 **2.0.9 Target Sequence Length (-x)**

64 The target sequence length is the length of a read expected using PCR. The forward and reverse primers
65 extract a portion of DNA between them. The value of this option is based on the primers you used, so
66 check their documentation or ask whoever did the PCR.

67 **2.0.10 Charybdis Scripts Directory (-g)**

68 Instead of setting the environment `$PATH` variable or placing the scripts in a systemwide bin directory,
69 the directory holding all the pipeline scripts is an option. Less conventional than `$PATH`-based methods,
70 but very easy to deal with.

71 **2.0.11 NCBI Taxonomy Database (-t)**

72 This database is used to get an organism's scientific name from a numeric taxonomic ID. Taxonomy
73 database downloaded from `ftp://ftp.ncbi.nlm.nih.gov/pub/taxonomy/`.
74 It is the same data as `https://www.ncbi.nlm.nih.gov/Taxonomy/Browser/wwwtax.cgi`

75 **3 PIPELINE ARCHITECTURE**

76 Pipelines are bash scripts that just string together a group of other scripts/programs to perform the metabar-
77 coding process. Example scripts are included in `/work/hobi/GCL/GCL_Charybdis/pipelines`.
78 The script called `charybdis_generic` is expected to handle most projects with little customization.

79 The section will describe all of the steps of the `charybdis_generic` pipeline. Hopefully others
80 will be able to use this information to create a pipeline tailored to other metabarcoding projects.

81 Throughout the text, *attribute* refers to a piece of metadata included in a FASTA sequence header.
82 These are in the form of key-value pairs. Many tools in the OBITOOLS package create and modify
83 attributes, so we made many of our own scripts follow this convention.

84 3.1 Merge Reads

85 **Script:** `/work/hobi/GCL/GCL_Charybdis/bin/mergeReads.sh`

86 **Slurm wrapper:** `/work/hobi/GCL/GCL_Charybdis/bin/mergeReads.slurm`

87 The forward and reverse reads are merged into a single FASTA file, and basic quality filtering is done.
88 The reads are matched to sample ID which are included as metadata in the FASTA sequence headers.

89 3.1.1 Split FASTQ files

90 In order to run the next sections in parallel, the forward and reverse read FASTQ files are split into equally
91 sized chunks using *fastq-splitter.pl*.

92 3.1.2 Join forward and reverse reads

93 After this point, the pipeline uses a single file of merged forward and reverse reads. This is done using
94 OBITOOL's *illuminapairedend* tool. We have it reject reads whose alignment score is under a
95 threshold of 40.

96 3.1.3 Convert FASTQ to FASTA

97 Convert the merged files to FASTA. This is done because many tools operate only on FASTA and because
98 we don't need the sequencer quality information after this point. We can always use the unique sequence
99 ID to access it from the origin FASTQ, if needed.

100 3.1.4 Remove unaligned sequences

101 We use OBITOOL's *obigrep* to select only those sequences whose alignment mode is "joined". *Obigrep*
102 can search for patterns of values of specified attributes.

103 3.1.5 Remove sequences whose alignment length is below threshold

104 We use OBITOOL's *obigrep* to select only those sequences whose alignment length is \geq a threshold. We
105 have this threshold set to 20.

106 3.1.6 Match sequences to sample ID based on barcode

107 OBITOOL's *ngsfilter* is used to attach an attribute called *sample* to each sequence. The match is allowed
108 to have 2 sequence errors by setting the flag `-e 2`.

109 3.1.7 Concatenate into single FASTA

110 The small chunks are concatenated into a single FASTA file.

111 3.2 Filter Reads

112 **Script:** `/work/hobi/GCL/GCL_Charybdis/bin/ObiToolsPipeline.sh`

113 **Slurm wrapper:** `/work/hobi/GCL/GCL_Charybdis/bin/filterReads.slurm`

114 This was the first script written. At the time, the size of the project was not known and this file was
115 intended to be the entire pipeline. Thus, the script is called *ObiToolsPipeline.sh* instead of the
116 more appropriate *filterReads.sh*.

117 3.2.1 Split into smaller files by sample ID

118 As before, we form a smaller number of files for parallel processing. However, the sample ID is used to
119 split files instead of an arbitrary size. The sample ID is added to the FASTA filename to keep track.

120 3.2.2 Remove unneeded attributes

121 OBITOOLS added a huge number of attributes in a previous step. We get rid of the irrelevant ones with
122 *obiannotate*. Note that attributes are never truly discarded since the sequence ID can always be used
123 to query the intermediate FASTAs generated at each step of the pipeline.

124 3.2.3 Keep only unique sequences

125 In order to lower the filesize (often substantially), duplicated sequences are merged into a single represen-
126 tative sequence. The *obimerge* attribute is used to record how many sequences are represented by that
127 sequence. This is done using *obiuniq*.

128 3.2.4 Remove PCR errors

129 PCR errors are filtered using `obiclean`. After numerous experiments, we found that the following
 130 options were performing well for a set of COI sequences from Gulf of Mexico fish. Results from projects
 131 where we reused these settings have appeared to be fine.

Table 2. Settings for `obiclean`

-r	0.5	Controls when less abundant sequences are labeled as variants of a similar more abundant
-d	1	Max number of differences to be considered variants
-H		Keep only a representative of a set of variants.

132 3.2.5 Remove low read length

133 The amplified sequence should be a certain size. Using our COI primers, we expect the size to be 313
 134 base pairs. However, the actual size is often a bit lower or higher. You supply your base pairs length as a
 135 pipeline argument. The allowed range is that size ± 15 .

136 3.2.6 Remove chimeras

137 VSEARCH is used to remove chimeras using the *de novo* method.

138 3.2.7 Concatenate into single FASTA

139 The small chunks are concatenated into a single FASTA file.

140 3.3 Collapse into OTUs

141 **Script:** `/work/hobi/GCL/GCLCharybdis/bin/ClusterOTU.sh`

142 **Slurm wrapper:** `/work/hobi/GCL/GCLCharybdis/bin/ClusterOTU.slurm`

143 3.3.1 Cluster with CROP

144 CROP clusters sequences into OTUs (Hao et al., 2011). CROP has a number of parameters that control
 145 the OTU assignment. Assigning these parameters is tricky, but the CROP manual has a recommended
 146 process for determining the values (TingChenLab, 2017). We initially used the recommendation, but then
 147 modified it to get better results. **However, we don't currently remember how we figured out some of the**
 148 **values used in our process.** You can check the manual and compare to ours, but note that `-b` and `-z` should
 149 only affect performance, not clustering. However, we set `-e` based on `-z` (following the manual), which
 150 does affect results. **We need to investigate further.**

Table 3. Settings for CROP.

NUMSEQS is the total number of sequences. *BP* is the target sequence length.

-r	5	Clusters of size <i>leqr</i> are considered rare
-s		Specifies that 97% similarity needed to cluster
-b	$NUMSEQS/50$	Number of blocks
-z	$(150000/BP) - 0.1 * (150000/BP)$	Max number of sequence a block can hold
-e	$[-z] * 10$	Number of MCMC iterations

151 3.3.2 Correct the OTU size

152 CROP records what sequences end up in an OTU, as well as the total number of sequences represented
 153 by an OTU. However, CROP has no knowledge that `obiclean` and `obimerge` already collapsed
 154 duplicate/similar sequences into. We call `CROP_size_fix.sh` to get the actual OTU counts.

155 3.4 Taxonomic Assignment

156 Taxonomic assignment can be done using either BLAST or VSEARCH. (SAP coming soon). The exact
 157 scripts and slurm wrappers depend on method selected.

158 3.4.1 Run taxonomic assignment program in parallel

159 Split sequences into chunks and BLAST in parallel. Uses the BLAST ignore list (provided in parameters)
 160 to skip assignment to unwanted sequences. For example, we typically try to avoid assignment to
 161 environmental sequences.

162 3.4.2 Concatenate into single FASTA

163 The small chunks are concatenated into a single FASTA file.

164 3.4.3 Convert to Charon format

165 Charon is the name of our own CSV format for assigned sequences. The purpose is that we can write
166 conversion scripts to have assignments from various taxonomic assignment programs into a standard. The
167 Charon file has columns specified by Table 4.

Table 4. Columns for Charon CSV file

Query Sequence ID	Assignment Sequence ID	Scientific Name	NCBI GI	Number of sequences in OTU
-------------------	------------------------	-----------------	---------	----------------------------

168 3.5 OTUs VS Samples

169 Formerly: *Critters VS Tubes*.

170
171 Will explain soon, but basically just a script that makes a column for each sample and a row for each
172 OTU. The cells have a count of the number of sequences of that OTU in that sample. Also has columns
173 of the scientific names for various phylogenetic levels. Also has extra information columns such as the
174 sequence, BLAST scores, etc. No limit on added extra columns to this CSV file.

175 4 EXAMPLE RUN

176 4.1 Prepare Data

177 4.2 Run Pipeline Script

178 The following bash script executes `charybdis_generic.sh`, an implementation of the GCL pipeline.
179 The code segment shown is a single line of bash, using `\` for clarity to break it into one line per option.

```
180 bash /work/hobi/GCL/GCL_charybdis/pipelines/charybdis_generic.sh \
181     -p Simons-H1.ATTACTCG-TAATCTTA.L001 \
182     -i /work/hobi/GCL/20180330-Simons/Simons-H1/in \
183     -o /work/hobi/GCL/20180330-Simons/Simons-H1/out \
184     -n 20 \
185     -b /work/hobi/GCL/db/NCBI.BLAST.DBs/nt \
186     -d /work/hobi/GCL/db/GI_lists/environmental.NCBI.nucl__SORTED.gi \
187     -c /work/hobi/GCL/db/BOLD.FULL.fasta \
188     -x 313 \
189     -g /work/hobi/GCL/GCL_charybdis/bin \
190     -t /work/hobi/GCL/db/TAXO
```

191 Each option will be explained, and the first 10 lines of each file will be displayed. This way, you can
192 compare the format of your input files to those used for the example. Do not trust the spacing seen here.
193 Tabs may become spaces, etc. See the official file format specifications.

194 **-p** Name of project.

195 We used `Simons-H1.ATTACTCG-TAATCTTA.L001` to match the names of the input FASTQ files. The important
196 thing is that each project should have a unique name for keeping projects separate/organized.

197 **-i** Input directory.

198 This is where the input files should be placed. A slight abuse of terminology, since `<projectname>.samples.txt`
199 is generated by the pipeline, by reading sample IDs from `<projectname>.barcodes.txt`. Note that the
200 FASTQ sequences were shortened to fit on the page. A sequence of five periods (`.....`) represents a removed
201 segment.

```
202 [ekrell@hpcm Simons-H1]$ ls in
203     Simons-H1.ATTACTCG-TAATCTTA.L001.barcodes.txt
204     Simons-H1.ATTACTCG-TAATCTTA.L001.forward.fastq
205     Simons-H1.ATTACTCG-TAATCTTA.L001.reverse.fastq
206     Simons-H1.ATTACTCG-TAATCTTA.L001.samples.txt
207
208 [ekrell@hpcm Simons-H1]$ head in/Simons-H1.ATTACTCG-TAATCTTA.L001.forward.fastq
209     @MISEQ01:242:000000000-BLT5R:1:1101:15746:1350 1:N:0:ATTACTCGTAATCTTA
```

```

210 ACTTGATAGACCTCGGGGTGGCCGAAGAACCAGAATAGGTGTTGGTAAAGAATTGGGTCTCC . . . .
211 +
212 CCCCCFFFFFGGGGGGGGGHGGGGGHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHH . . . .
213 @MISEQ01:242:000000000-BLT5R:1:1101:16137:1376 1:N:0:ATTACTCGTAATCTTA
214 ATCACGGGAACCTGGATGAACAGTTTATCCTCCCTTGCCGGCAACCTGGCCCACGCAGGGGC . . . .
215 +
216 CDDDDDBCCBCFGGGGGGGGGHGGGGGHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHH . . . .
217 @MISEQ01:242:000000000-BLT5R:1:1101:15575:1386 1:N:0:ATTACTCGTAATCTTA
218 GTCCGCGTACTGGATGAACAGTATACCCCTTTAGCAGCAGCCATTGCACATGCAGGTGC . . . .
219
220 [ekrell@hpcm Simons-H1]$ head in/Simons-H1.ATTACTCG-TAATCTTA.L001.forward.fastq
221 @MISEQ01:242:000000000-BLT5R:1:1101:15746:1350 1:N:0:ATTACTCGTAATCTTA
222 ACTTGATAGACCTCGGGGTGGCCGAAGAACCAGAATAGGTGTTGGTAAAGAATTGGGTCTCCCCA . . . .
223 +
224 CCCCCFFFFFGGGGGGGGGHGGGGGHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHH . . . .
225 @MISEQ01:242:000000000-BLT5R:1:1101:16137:1376 1:N:0:ATTACTCGTAATCTTA
226 ATCACGGGAACCTGGATGAACAGTTTATCCTCCCTTGCCGGCAACCTGGCCCACGCAGGGGCATCA . . . .
227 +
228 CDDDDDBCCBCFGGGGGGGGGHGGGGGHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHHH . . . .
229 @MISEQ01:242:000000000-BLT5R:1:1101:15575:1386 1:N:0:ATTACTCGTAATCTTA
230 GTCCGCGTACTGGATGAACAGTATACCCCTTTAGCAGCAGCCATTGCACATGCAGGTGCATCT . . . .
231
232 [ekrell@hpcm Simons-H1]$ head in/Simons-H1.ATTACTCG-TAATCTTA.L001.barcodes.txt -n 5
233 #exp sample tags forward_primer reverse_primer
234 Simons-H1 1901-B ACAGTG:CCGTCC GGWACWGGWIGA . . . . TANACYTCNGGRTGN . . . .
235 Simons-H1 1685-C ACAGTG:GTTTCG GGWACWGGWIGA . . . . TANACYTCNGGRTGN . . . .
236 Simons-H1 1759-B ACTGAT:ATGTCA GGWACWGGWIGA . . . . TANACYTCNGGRTGN . . . .
237 Simons-H1 1774-X ACTGAT:GCCAAT GGWACWGGWIGA . . . . TANACYTCNGGRTGN . . . .
238
239 [ekrell@hpcm Simons-H1]$ head in/Simons-H1.ATTACTCG-TAATCTTA.L001.sampledescs.csv
240 -n5
241 Tube,Description
242 01-A, Larimus_fasciatus
243 05-B, Larimus_fasciatus
244 05-C, Larimus_fasciatus
245 08-A, Larimus_fasciatus
246
247 -o Output directory.
248 Files generated by the pipeline, including numerous intermediate file, are placed here.
249 -n Number of parallel instances.
250 The nodes on TAMUCC's HPC cave 20 cores, so 20 were used.
251 -b BLAST Database.
252 Specify path to BLAST database. We are using the nucleotide (nt) database,
253 which is downloaded from ftp://ftp.ncbi.nlm.nih.gov/blast/db/.
254 -d List of GIs for BLAST to ignore This is a list of GIs, where each GI specifies a sequence in the BLAST
255 database to ignore. This file is mandatory, but may be blank (See Section 5.2).
256 We attempt to remove environmental DNA, so our list has GIs of environmental sequences.
257
258 [ekrell@hpcm Simons-H1]$ head -n 5 . . . . /environmental.NCBI_nucl_SORTED.gi
259 1000014437
260 1000014439
261 1000014441
262 1000014443
263 1000014445
264
265 -c Chimera Database. We use a FASTA file generated from the BOLD database. But, we ended up doing de novo
266 chimera detection. Currently this file is required and just not being used (See section 5.3).
267 -x Target sequence length (basepairs). This number is based on the primers used during PCR.
268 -g Directory of charybdis scripts.
269
270 ls /work/hobi/GCL/GCL_charybdis/bin
271 \ fastq-splitter.pl

```

```

268 AddBlast.slurm filterReads.slurm
269 AddVsearch.slurm mergeReads.sh
270 AssignToMarkers.sh mergeReads.slurm
271 AssignToMarkers.slurm ObiToolsPipeline.sh
272 blast_10custom_to_charon_OTU.R ObiToolsPipeline.slurm
273 BlastMeta.sh ObiToolsPipeline_stats.sh
274 BlastMeta.slurm OTU_CVT_addBlast.R
275 ClusterOTU.sh OTUvsTube.slurm
276 ClusterOTU.slurm split_by_marker.R
277 CombineAndCROP.sh vsearch_blast6custom_to_charon-BLASTDB_OTU.R
278 crittersVStubes_OTU.R vsearch_blast6custom_to_charon_OTU.R
279 CROP_size_fix.sh vsearch_getTAXIDfromBOLDseqid.sh
280 CROP.slurm Vsearch.sh
281 determine_marker.R Vsearch.slurm
282 determine_marker.sh

283 -t Directory with NCBI taxonomy database Taxonomy database downloaded from ftp://ftp.ncbi.nlm.
284 nih.gov/pub/taxonomy/.
285 -g Directory of charybdis scripts.

286 [ekrell@hpcm pipelines]$ ls /work/hobi/GCL/db/TAXO
287 Accession2Taxid delnodes.dmp gc.prt merged.dmp nodes.dmp taxdump.tar.gz
288 citations.dmp division.dmp gencode.dmp names.dmp readme.txt

```

289 5 AREAS FOR IMPROVEMENT

290 5.1 Support NCBI Accession IDs

291 NCBI is phasing out GIs (sequence IDs), which this pipeline uses at various steps (NCBI, 2016). Accession IDs are
292 the current standard.

293 5.2 Pipeline should not require BLAST ignore list

294 Currently this is required. If you don't need it, you can supply a blank file. Not very elegant. Would be simple to
295 make it optional.

296 5.3 Chimera database should not be mandatory

297 The generic pipeline uses *de novo* chimera detection, but still requires the chimera database. A better solution would
298 be to default to *de novo*, but use the database if supplied.

299 6 STATISTICAL ASSIGNMENT PACKAGE (SAP)

300 **We have SAP working, but not meshed into the pipeline yet. So it has its own section, at least for now.**

301 SAP (Munch et al., 2008) differs from typical assignment software (BLAST, VSEARCH) in that it is based on
302 computing a posterior probability that the query sequence is a member of a particular clade. Two major phases are
303 involved in an SAP taxonomic assignment. First, BLAST is executed to find a set of similar sequences. Since similar
304 sequences should indicate similar biological properties, the set of accepted BLAST sequences are called the set of
305 homologues with the query. To strengthen the assumption that the sequences are homologous, there is a limit on
306 how distant the furthest apart homologues can be. Several parameters are used to control BLAST's sequence search
307 and also to control how SAP chooses which of those sequences to include in the homologue set. It can and does
308 happen that a single species has several sequences in the database. However, a particular taxonomic group appears
309 only once in the phylogenetic tree. Therefore, SAP focuses on diversity and only uses the best match for a given
310 species. While this is necessary to construct meaningful trees, it also means that the species which only appears once
311 in a set of BLAST results is given equal weight to a predominate species result. The set of homologues and the query
312 sequence are used to sample a large number of trees using Markov Chain Monte Carlo. The posterior probability for
313 any taxonomic group is based on the proportion of trees in which the query sequence was in that group. A simple
314 case that shows how this can differ from BLAST occurs when several species from the same genus are present in the
315 homologue set. Where BLAST would simply list the hits in order of sequence ID, SAP determines that it cannot
316 reliably be determined what species the query belongs to. A top hit may differ little from the hits just below it, and it
317 is essentially arbitrary to assume that the top hit is the exact species when it scored closely against other species as
318 well. The SAP result would give a low posterior probability to any particular sequence, but the genus would have a
319 very high probability.

320 *The above text taken from GCL's work-in-progress manuscript "An Evaluation of Software for Taxonomic*
321 *Assignment with DNA Metabarcoding"*

322	6.1 Fixing SAP
323	6.2 Running SAP
324	6.3 Evaluating SAP
325	6.3.1 Categorize SAP results
326	6.3.2 Generate SAP report
327	6.4 SAP outstanding issues
328	6.5 Complete SAP example
329	REFERENCES

- 330 Hao, X., Jiang, R., and Chen, T. (2011). Clustering 16s rRNA for OTU prediction: a method of unsupervised bayesian
331 clustering. *v*, 27:611–8.
- 332 Munch, K., Boomsma, W., Huelsenbeck, J. P., Willerslev, E., and Nielsen, R. (2008). Statistical assignment of DNA
333 sequences using bayesian phylogenetics. *Systematic biology*, 57(5):750–757.
- 334 NCBI (2016). Ncbi is phasing out sequence GIS – here’s what you need to know.
- 335 Ratnasingham, S. and Hebert, P. D. (2007). bold: The Barcode of Life Data System (<http://www.barcodinglife.org>).
336 *Mol. Ecol. Notes*, 7(3):355–364.
- 337 Rognes, T., Flouri, T., Nichols, B., Quince, C., and Mahe, F. (2016). VSEARCH: a versatile open source tool for
338 metagenomics. *PeerJ*, 4:e2584.
- 339 TingChenLab (2017). Crop. <https://github.com/tingchenlab/CROP>.